

Illinois cases linguistic analysis

Davide Carletti and Francesco Tomaselli

University of Milan

Abstract. Court judgments temporal analysis

Keywords: Legal texts · Topic modelling · Word embeddings

1 Introduction

The goal of this project is to study words among the temporal axis in court decisions. More precisely, we aim to find information about relevance and frequency of single or multiple terms, and also to discover correlations between them.

2 Research question and methodology

The goal of the project is to analyze, from a collection of court decisions, the relevance of terms in time while also gathering information about possible correlations between them. The reasons to perform such a study are various, firstly, one could find interesting and unexpected correlations between terms, also, studying them in time could reveal a change of context of some words.

The research question is focused in particular on three main areas of interest: narcotics, weapons and investigations. Thus, words related to this three main sets are the most interesting ones to analyze.

In the first place, the methodology aimed to separate the original dataset in three main areas, narcotics - weapons - investigations indeed; in this way, each document would have been assigned to one area. Having documents divided per area of interest, it would have been possible to divide documents in temporal periods and perform topic modelling and language modelling on the temporal subsections of the three areas. This analysis would have allowed to directly confront the three sets also in a temporal way, showing differences in the language and topics between words regarding drugs, narcotics and investigations respectively. Some scientific literature has been explored and tested, based on the idea of Guided topic modelling (paragraph 3.2 - Topic Modelling), but it has been soon discovered that the dataset comprehended many other topics other than the three main areas of interest, and most importantly even when finding documents and terms regarding them, it was pretty difficult to separate them in three distinct sets, because in reality the three sets are interconnecting each other, being usually used together.

After this discovery, the methodology has been changed to a more general one: the dataset has been explored, by finding the various topics and language

models composing it, finding also semantic shift in words context between different years and periods of time. After finding more generic topics, the discovered topics have been analyzed (and labelled), and the ones regarding the initial three areas of interest have been chosen to be explored more in depth, to find more meaningful results for the original research question.

3 Experimental results

3.1 Dataset preprocessing

The dataset in use is the Illinois Bulk Dataset, that contains 183146 cases with 194366 judges opinions.

The first step of the preprocessing is to merge the opinions about a case into one, obtaining a single document for each dataset entry. Then, each document goes through a text cleaning and tokenization phase, where the first part is done with the help of regular expressions, while the second uses Spacy to obtain a list of terms.

3.2 Topic modelling

To have an overview of the topics discussed on the dataset a Latent Dirichlet Allocation model is trained on the tokenized texts. One of the key parameters of such a model is the number of topics, and, given the fact that cases could potentially talk about anything, an *Halving search* is performed to find a good value. [?]

We opt for an halving search since the number of topics could be anything, we fixed a range between ten and thirty and training each model to then evaluate results would take a huge amount of time. Halving search mitigates the problem, as it trains firstly on smaller datasets, select the best models, and retrain on bigger slices of data until a final model is found. This methodology can be ten times faster than grid search. The selection criterium for the search is the Log Likelihood of each model. The search revealed that the optimal number of topics is 14.

The resulting topics are promising but the words of interest, namely a collection of narcotics, investigation and weapons terms, are merged together in a few topics. To solve the issue we decide to run topic modelling on a subset of the previous topics with the same technique as before. The result is similar, we find again 14 topics, but this time they are much more specific.

3.3 Temporal word embeddings

One of the objective of the project is to find correlations among words and one tool that can be used is word embeddings. The technique assigns a real vector of a given dimensions to each word in the document collection, creating a way to directly compare the context similarity between two terms.

For this task we used Gensim's Word2Vec implementation and trained different models in three ways:

1. *Global model*: trained on the entire document collection, with 100 components vectors;
2. *One year models*: they are trained on subsets of the dataset divided by years;
3. *Ten years models*: trained on epochs of 10 years each.

The first model gives information about the whole dataset, it can be used to compute similar words queries, while the others can be exploited to find context and semantic shifts among the temporal axis. Taking inspiration from the Hist-Words [1] work on semantic shift, we start by aligning the models, and then find, given a word and a base year, the similarity of that word in time with respect to the base year. This approach can reveal if a word changed semantic or context, and when, with respect to a given year, an example can be seen in 1.

Similarly, the ten year models are aligned, but this time, given a term, we compute the difference among two consecutive epochs. The idea is similar to the previous one but slightly different, this time a drop in the similarity sequence reveals a change of meaning from an epoch to the other, the resulting sequence can be seen in 2 on the next page.




Fig. 1. Semantic shift of the word assault with respect to the year 2010, we can see a drop if similarity around 1960.

3.4 Webapp

In order to visualize and explore the results of the analysis, a web interface has been developed. Through the UI, the user can search multiple words united (separated by ",") and/or confront words (separated by "-"); the app will show various interesting sections:

- Top 15 words of similar context of the searched words (using word embeddings).
- Frequency graph of documents in time containing the searched words.




Fig. 2. Similarity between consecutive epochs of the term assault, we can see a drop between 1980-1970 and 1970-1960.

- Semantic shifts of the searched words, both by epoch (two adjacent epochs of 10 years) and by single year. If the user clicks on a point of the graph, the contexts of the selected word in the different selected years/epochs will be showed, allowing the comparison of contexts which helps understanding the evolution of words meaning and usage during time.
- Topic distribution of the searched words (both generic and areas of interest-specific topics). By clicking on the showed radar chart, or on the lateral list of topics, the selected topic info will be loaded in the below section.
- Most important words of the selected topic. There are three ways of visualization: a barchart with the first 15 words, a Wordcloud with the first 80 and a Treemap with the first 50.
- Topic distribution over time and over the different Illinois courts.

The application has been developed in Python and is accessible at <https://illinois-cases-analysis-webapp-qka7d4ktba-ew.a.run.app/>.

3.4.1 Webapp developing technology

3.4.2 Cloud hosting

4 Concluding remarks

References

1. William L. Hamilton, Jure Leskovec, and Dan Jurafsky. ACL 2016. Diachronic Word Embeddings Reveal Statistical Laws of Semantic Change.